

ProvLoom: Security Explanation of Agentic Skills via Runtime and Instruction Provenance Chains

Anonymous Authors
Anonymous Institution

Abstract—Agentic skills integrate local files, tool invocations, network actions, scripts, configuration documents, maintenance instructions, and LLM-mediated decision-making into complex execution workflows. While existing security analysis tools can flag potentially risky behaviors, they often fail to determine whether observed evidence forms a complete source-to-sink attack chain, and they may overlook risks embedded in setup or maintenance instructions.

In this paper, we present ProvLoom, a security explanation framework for agentic skills based on runtime and instruction-derived provenance chains. ProvLoom decouples dynamic execution from static instruction analysis. The runtime component constructs an execution provenance graph from sandbox telemetry, while the instruction component analyzes local artifacts to identify trust boundary violations, control transfers, and security-relevant dependencies. These two views are unified into a provenance-based reasoning model that classifies each case as a runtime chain, instruction-derived chain, hybrid chain, or an incomplete (open) chain, enabling explicit reasoning about whether a closed source-to-sink path exists. We build Benchmark, a controlled dataset of 147 agentic skills, including 104 malicious and 43 benign cases, to evaluate chain reconstruction and benign-hard discrimination. On executable cases, ProvLoom achieves 0.9765 endpoint accuracy, 0.9338 edge F1 score, and 0.8676 complete chain recovery rate. Compared with existing security tools for agentic systems, ProvLoom uniquely determines whether evidence forms a closed provenance chain and provides structured, actionable explanations. We further conduct a real-world security audit of the public agentic skill ecosystem. We collect 88,769 skills from skill markets and apply automated prescreening to prioritize 670 samples for manual analysis, identifying 21 malicious, 79 ambiguous, and 570 benign cases. ProvLoom successfully reconstructs closed provenance chains for 9 of the 21 confirmed malicious samples, including both newly discovered and previously reported cases. These results demonstrate that ProvLoom enables evidence-based triage and improves the interpretability of security assessment in large-scale agentic skill ecosystems.

Index Terms—Skill Security, Supply Chain Security, Dynamic Analysis, Vulnerability Detection

1. Introduction

Large language model agents connect model reasoning

to external tools and local files to transform textual outputs into system behaviors [2], [36], [38], [51], [52], [54]. Tool invocation translates model decisions into concrete actions where harmful outputs may trigger unintended file operations or network requests [13], [22], [39], [52]. Reusable skill packages bundle task instructions with helper scripts and metadata to facilitate capability reuse across diverse tasks [5], [21], [45], [53], [55].

Risky behaviors within these packages can originate from local instructions and propagate through helper code or setup steps to reach sensitive system endpoints [2], [10], [46], [52], [54]. Malicious skills may embed unsafe payloads across multiple components to steer agents toward unauthorized data access or command execution via prompt injection and tool misuse [5], [10], [13], [21], [22], [39], [44], [52], [53]. Reports indicate that various agent platforms have hosted skill packages mimicking legitimate functionality while potentially stealing credentials or hijacking tool chains [13], [22], [39], [52]. These findings suggest that skill packages warrant systematic detection and evaluation as potential attack vectors in agent ecosystems [5], [11], [21], [29], [35], [41], [53], [55], [56].

Prior Studies. Prior work on skill security studies risks introduced by reusable skills. Existing studies report malicious instructions, data exfiltration, unsafe command construction, overbroad capability use, risky helper scripts, and unsafe setup or maintenance behavior [5], [21], [53]. Agent security work studies how untrusted content affects tool selection, file access, command execution, and network behavior when agents act through tools [2], [10], [13], [22], [39], [44], [46], [52]. Software supply chain work shows that installers, package managers, update channels, and maintainer controlled setup paths can introduce risks before normal application execution [11], [29], [35], [41], [56]. Prior work therefore shows that skill risks can arise during agent execution and during package distribution.

Recent skill scanners and agent security tools audit the risks by inspecting skill packages and agent application files [1], [5], [8], [21], [37], [47], [53]. The tools commonly examine local artifacts, helper code, dependency files, commands, URLs, capability declarations, and setup text. Outputs usually take the form of rule hits, suspicious indicators, severity labels, static findings, or reachability verdicts. Scanner outputs support screening and help reviewers reduce a corpus to a smaller set for manual review.

However, static scanner outputs still leave practical re-

view gaps. Most tools analyze local artifacts and do not show whether a suspicious command, URL, permission, or data access occurs during execution. Risk labels can hide different evidence states, and some tools do not provide severity levels that explain where a reported risk comes from. Benign documentation, defensive examples, normal setup text, and public API usage can match patterns found in malicious skills. Static screening can also miss behavior that depends on generated files, tool calls, model selected actions, credentials, external state, or runtime branches.

Our Study. In this study, we propose ProvLoom, a dynamic framework that explains skill security decisions through provenance chains with explicit evidence types. The framework checks whether runtime evidence closes a chain from a source, through relays, to a sink. The instruction analysis checks whether setup or maintenance instructions close a path that crosses a trust boundary, transfers control, and leads to a security impact. The final explanation classifies each case as a runtime chain, a chain derived from instructions, a hybrid chain, or no closed chain. ProvLoom complements broad scanners, host provenance systems, and static analyzers. ProvLoom provides users with a compact explanation of whether a skill should be escalated and what evidence supports the decision.

We build Benchmark to evaluate whether a skill analyzer can recover security chains. Benchmark contains 147 skill cases. The malicious cases cover six attack families. The benign cases are designed as hard controls. Some benign cases contain outward looking wording, commands, URLs, or setup text but do not form malicious chains. Each case is annotated before evaluation with the expected source, relay, sink, root cause, and runnability boundary. We evaluate ProvLoom by comparing the recovered chains with the annotated chains in Benchmark. The comparison shows that runtime provenance gives users more complete evidence than behavior detection alone.

The public corpus audit applies ProvLoom to more than 80,000 public skills from skill markets. We manually annotate 670 selected samples from this audit corpus and identify 21 malicious samples, 79 ambiguous samples, and 570 benign samples. ProvLoom recovers closed chains for 9 of the 21 malicious samples in the annotated subset. The confirmed malicious cases often involve OpenClaw agent and CLI supply chain behavior, including protected archives, remote script execution, and setup or maintenance control transfer. The audit shows how closed chain evidence supports triage by separating confirmed malicious cases from ambiguous cases and benign cases with only indicators.

Contributions. We make the following contributions:

- We propose ProvLoom, a dynamic framework for skill security explanation. ProvLoom converts sandbox telemetry into an execution provenance graph to recover runtime source to sink chains. The instruction analysis component examines setup and maintenance artifacts to recover chains derived from instructions. The decision layer assigns an evidence type only when the relevant path closes.

- We build Benchmark, a suite of 147 cases for chain recovery and hard benign discrimination. Benchmark contains 104 malicious and 43 benign cases. The malicious cases cover exfiltration, unauthorized transfer, unsafe command construction, LLM mediated unsafe actions, and mixed multistep flows.
- We evaluate ProvLoom and baseline analyzers on Benchmark. ProvLoom reaches 0.9765 endpoint accuracy, 0.9338 edge F1, and 0.8676 complete chain rate on the completed dynamic cases.
- We audit 88,769 skills from the real world. ProvLoom identifies 21 malicious samples and recovers closed chains for 9 confirmed malicious samples. We have conducted responsible disclosure for confirmed samples.

2. Background

2.1. Agentic Skills

Large language model agents can use external tools, local files, scripts, network endpoints, and system commands during task execution [2], [36], [38], [51], [52], [54]. A skill is a reusable unit that packages instructions and implementation artifacts for an agent to invoke [45], [55]. A skill bundle may include local instructions, README files, `SKILL.md`, package metadata, tool descriptions, helper scripts, and other files that describe how the skill should be installed, configured, and used [5], [21], [53]. Figure 1 summarizes a typical skill usage flow.

A skill bundle can contain both natural language guidance and executable artifacts [21], [55]. The natural language files describe the task, usage conditions, inputs, and expected outputs. The implementation artifacts provide helper code, command templates, tool descriptions, dependency information, or configuration steps [5], [21]. An agent can use these artifacts to decide when the skill applies and how to carry out the requested task [45], [55].

Skill execution usually involves several stages. The agent first reads the skill description and selects a skill that matches the user request [45], [55]. The agent then follows the local instructions, calls declared tools when needed, and may invoke helper scripts or access files described by the bundle [36], [51]. Setup and maintenance documents may also describe installation, configuration, update, or account connection steps [5], [21], [47]. Skill behavior can come from both execution artifacts and local instruction artifacts.

2.2. Static and Dynamic Analysis

Static analysis examines artifacts before execution. In the skill setting, analyzed artifacts may include local instructions, helper scripts, metadata, dependency declarations, command strings, URLs, and setup or maintenance text. Static analysis can support screening and can identify source to sink paths before a skill runs [23], [40], [42], [50].



Figure 1. Typical skill usage flow.

Dynamic analysis examines concrete execution [7], [14], [16], [33]. A dynamic run can collect file operations, process creation, command execution, network actions and tool calls [7], [16], [33], [46]. Runtime provenance systems use relations among runtime entities to explain how an observed outcome occurred [7], [14], [16], [33]. A dynamic report describes runtime evidence under a bounded input, configuration, and environment [10], [46], [53]. A bounded run may miss behavior that requires credentials, external platform state, user interaction, delayed triggers, rare inputs, or nondeterministic choices [10], [21], [53].

2.3. Directed Graphs and Provenance Paths

A directed graph represents objects as vertices and ordered relations as edges. Provenance models use directed graph structure to describe how entities, activities, and results depend on earlier entities and activities [28]. Security provenance systems use similar graphs for runtime analysis. Vertices may represent processes, files, sockets, network endpoints, tools, or other runtime objects, and edges may represent operations or information flow [4], [16], [33].

Graph paths help security analysis move beyond events. A file read, a file write, and a network connection may appear in the same trace without belonging to the same behavior. A path shows whether observed relations connect those events and exposes relays such as generated files, helper processes, or tool calls [7], [14], [16], [33].

Source, relay, and sink provide a compact vocabulary for security paths. A source is where data or control originates. A sink is an endpoint or operation where data leaves the system, changes persistent state, or causes security impact. A relay is an intermediate entity that connects a source and a sink. A closed chain exists only when recovered

relations connect the source, relays, and sink in one path. Disconnected indicators remain weak evidence.

3. ProvLoom Detection Framework

3.1. ProvLoom Overview

ProvLoom audits an agent skill by checking whether the available evidence forms a closed security chain. The input contains a skill bundle and an execution context. The skill bundle includes local artifacts such as `SKILL.md`, helper scripts, metadata, and configuration files. The execution context defines the bounded run used to collect dynamic observations. The output is a compact review report with an evidence type, a recovered chain, a root cause, a severity level, and traceability to the supporting evidence. Figure 2 summarizes this workflow from skill inputs to evidence analysis and report generation.

Given a skill bundle, ProvLoom constructs two evidence surfaces. Dynamic evidence comes from bounded execution. Static instruction evidence comes from local skill artifacts. These two evidence surfaces are analyzed separately so that observed behavior and latent setup or maintenance behavior do not get mixed during chain recovery.

During bounded execution, ProvLoom runs the skill in the provided execution context and collects telemetry from the run. The telemetry records file operations, process events, network actions, tool calls, and decisions mediated by the LLM. Dynamic analysis converts these observations into an execution provenance graph. It reports an observed chain only when the graph connects a source to a sink through observed relations. Disconnected dynamic observations remain weak signals.

In parallel, static instruction analysis scans local artifacts in the skill bundle. The analysis looks for setup or maintenance text that describes external acquisition, local execution, persistence, global state modification, bulk update behavior, or sensitive account connection. It reports a latent chain only when the artifacts support a path from trust boundary crossing to control transfer and security impact.

After the two analyses finish, the attribution layer compares the dynamic result and the static instruction result. The layer assigns the evidence type without merging the two evidence surfaces. The report records whether the chain is supported by dynamic evidence, static instruction evidence, or both. It records the recovered chain, the root cause, a severity level from critical, high, medium, and low, and the events or local artifacts that support the decision.

3.2. Threat Model

We define the audit setting along three axes: the skill artifacts under review, the adversary who can supply or modify them, and the audit goal that ProvLoom pursues.

Targets. ProvLoom audits agent skill bundles whose security behavior may be defined across local artifacts. A bundle can contain natural language instructions, tool descriptions, command templates, helper scripts, dependency




Figure 2. Overview of the ProvLoom analysis pipeline.

declarations, URLs, setup text, maintenance text, metadata, and configuration files. These artifacts may affect skill use, installation, update, or maintenance.

Adversary capabilities. We consider an attacker who can publish a malicious skill bundle or modify an existing bundle. The adversary controls the artifacts under review, including instructions, tool descriptions, command templates, helper scripts, dependency declarations, URLs, setup text, and maintenance text. The goal is to trigger risky behavior during skill use, installation, update, or maintenance, such as sensitive data access, unsafe command execution, external code installation, or persistent state changes. The adversary does not directly control the user’s agent, host, credentials, or external platform state. Behavior outside the analyzed execution is in scope only when the skill triggers it during the bounded run or when local artifacts describe a setup or maintenance path that a user or maintainer may follow.

Audit Goal. ProvLoom aims to recover closed security chains from dynamic evidence and static instruction evidence. Static screening can flag a URL, command template, permission request, archive operation, or other suspicious artifact, but these indicators do not show whether the setup path reaches a security impact. A bounded run can provide stronger evidence by connecting observed file access, tool calls, command execution, network transfer, or state changes into a source, relay, and sink chain. At the same time, bounded execution may miss risks that require user confirmation, credentials, external platform state, or later maintenance steps. ProvLoom uses static instruction evidence for these latent setup and maintenance paths, and it separates

such risks from ordinary documentation, defensive examples, and incomplete indicators. The final report records the recovered chain, root cause, evidence type, severity level, and supporting traceability.

4. The Design of ProvLoom

4.1. Local Artifact Loading

ProvLoom loads the local artifacts in a skill bundle. The loader parses SKILL.md, README variants, package metadata, local configuration files, declared action entries, setup documents, maintenance notes, and selected local scripts. The loader records which files define the skill interface, which files define helper behavior, and which files describe setup or maintenance actions. The loader does not execute setup instructions.

The loading stage gives analysis a stable artifact boundary. Dynamic analysis uses the loaded interface and helper action entries to decide which actions can run under bounded conditions. Instruction analysis uses the setup and maintenance files to recover latent paths described by artifacts.

4.2. Sandboxed Dynamic Analysis

ProvLoom executes selected helper actions and declared tool actions in a sandboxed execution context. The sandboxed dynamic analysis is a lightweight execution layer for skill review, not a full operating system monitor, skill platform, or host provenance substrate. Its scope is limited to

bounded runs, controlled tool interfaces, and event collection needed by ProvLoom.

The sandboxed execution context defines the workspace, synthetic sensitive files, time budget, resource limits, and available tool interfaces. Real credentials and unrestricted host access are outside this context. During a run, the execution mediator records action selection by the LLM when a tool call triggers local side effects. File observations cover reads and writes in the sandbox workspace and selected synthetic sensitive files. Process observations record command invocation, parent process relation, normalized arguments, and exit status when available. Network observations record destination host or URL class, request method when available, and tool linkage. Full packet contents are excluded from the collected evidence.

The telemetry exporter converts these observations into deterministic event records for dynamic chain recovery. Each record stores the execution order, event kind, acting entity, object entity, and event context. Event kinds include read, write, execute, connect, tool call, and decision. Event context stores the execution step, parent reference, normalized argument summary, status, and tool marker when available. These records provide the input to dynamic analysis and the evidence references used in review output.

4.3. Dynamic Chain Recovery

Dynamic chain recovery takes the event records produced by sandboxed dynamic analysis and reconstructs causal paths for behavior observed during the bounded run. The goal of this stage is to decide whether dynamic events form a closed chain from a source to a sink.

The normalizer converts dynamic events into an execution provenance graph. Vertices represent processes, files, network endpoints, tools, data objects, and LLM decisions. Edges represent reads, writes, connections, causal relations, and data flow relations. A file read creates a read edge from the acting process or tool to the file. A file write creates a write edge to the written file. A process invocation creates a causal edge from the parent to the child process. A network action creates a connection edge to the destination endpoint. A tool call links the LLM decision to the tool and connects the tool to its resulting file, process, or network effects. Each vertex and edge keeps a back reference to the event records that support it, so the final report can trace every recovered step to concrete dynamic evidence.

Dynamic traces often contain loader activity, interpreter state, temporary files, environment reads, and helper calls. The normalizer keeps events that can support a source, relay, sink, or causal context. It marks common noise with a noise tag during graph construction. Chain reconstruction uses this tag during path ranking and filtering.

The execution provenance graph separates behavior labels from causal explanation. A behavior label may indicate that a sensitive file was read and that an HTTP request occurred. The graph confirms a dynamic chain only when the sensitive source and external sink belong to the same connected causal path. It also records the tool call, helper

process, intermediate file, or generated artifact that links the source to the sink.

A dynamic chain connects a source, relays, and a sink through read, write, connect, causal, or data flow relations. Sources include sensitive files, credential artifacts, generated reports, and user provided data objects. Relays include intermediate files, helper processes, tool calls, and generated artifacts. Sinks include network endpoints, commands, persistence targets, and global state modifications when observed. ProvLoom reports a closed dynamic chain only when the graph contains a connected path from a source to a sink. Disconnected observations remain weak signals.

The chain reconstructor extracts a compact explanation from the graph. It ranks candidate paths by path length, number of noise nodes, source priority, sink priority, and a lexical tie breaker. Source ranking prioritizes sensitive files associated with tool calls and generated artifacts. Sink ranking prioritizes declared URL endpoints and HTTP requests associated with tool calls. The exported chain keeps the source, sink, and relays. After the primary path is selected, the filtering step removes loader, interpreter, and temporary nodes that do not explain the source to sink connection.

4.4. Static Instruction Chain Recovery

Static instruction chain recovery analyzes local skill artifacts to recover latent setup and maintenance risks. The goal is to determine whether local documents support a closed path from trust boundary crossing to control transfer and security impact.

The analyzer converts local artifacts into a typed local artifact graph. Vertices represent documents, extracted actions, external resources, execution targets, capability contexts, and impact sinks. Edges encode document supported relations such as mention, instruction, acquisition, execution, update, and capability grant. Each vertex and edge keeps a back reference to the local file span that supports it.

A closed instruction chain has three parts. A trust boundary crossing introduces content or capability into the setup or maintenance path, such as a remote script, downloaded binary, protected archive, or external account connection. A control transfer step gives that content or capability a way to affect skill behavior, such as local execution, an install hook, an updater command, or replacement of local skill files. A security impact sink captures the resulting change, such as persistence, global environment modification, broad update authority, or sensitive account access.

The scan stays within local artifacts. The analyzer reads files such as `SKILL.md`, `README` variants, setup documents, and local scripts. Remote content, installer execution, script execution, and dependency inspection are outside this stage. A setup path described in local artifacts remains latent until dynamic evidence observes execution.

Extracted indicators, such as external code acquisition, protected archives, persistence setup, bulk update behavior, and sensitive capability context, become strong static instruction evidence only when the local artifact graph closes the required path. Bare links, ordinary package installation,

standard API use, and defensive examples remain weak or suppressed evidence.

4.5. Evidence Attribution

The attribution layer turns the dynamic result and the static instruction result into one review outcome. ProvLoom promotes a finding only when dynamic events or local artifacts form a closed typed path. Isolated indicators, disconnected observations, and incomplete artifact paths remain supporting context. The attribution layer also assigns a severity level to each review outcome, using the labels critical, high, medium, and low.

The decision layer assigns one of four evidence types. *Observed evidence* means that bounded execution exposes a closed dynamic chain. *Instruction evidence* means that local artifacts expose a closed latent setup or maintenance chain. *Hybrid evidence* means that both evidence surfaces close chains. *No chain* means that ProvLoom finds only weak indicators, disconnected dynamic observations, or incomplete local artifact paths.

The evidence type defines how the report should be interpreted. Observed evidence describes behavior seen during bounded execution. Instruction evidence describes behavior prescribed by local setup or maintenance artifacts. Hybrid evidence keeps the dynamic chain as the primary explanation and attaches the instruction chain as latent supporting evidence. No chain preserves weak signals without presenting them as a complete causal explanation.

Attribution follows the closed path structure. When a closed dynamic path exists, ProvLoom uses that path as the primary explanation. A closed static instruction path can raise review priority when dynamic evidence is partial or unavailable. The static instruction path must still contain a trust boundary crossing, a control transfer step, and a security impact sink. ProvLoom derives the root cause from vertices, edges, and sink types in the closed path. Root cause labels include unsafe data flow, unsafe command construction, LLM mediated unsafe action, overprivileged tool use, supply chain setup, environment control setup, and unknown. The severity level is assigned from the impact sink, evidence type, sensitivity of the source, and privilege of the affected action.

Algorithm 1 records the attribution procedure used by ProvLoom. Dynamic evidence and static instruction evidence are tested separately. Closed paths become reviewable chains. Weak indicators stay below the chain threshold.

Algorithm 1: Evidence Type Chain Decision

Input: dynamic event records E_D , local instruction artifacts A_I
Output: severity level S , primary chain C , root cause R , evidence type T

- 1: $G_D \leftarrow$ build execution provenance graph from E_D
- 2: $C_D \leftarrow$ extract dynamic chain from G_D
- 3: $G_I \leftarrow$ build local artifact graph from A_I
- 4: $C_I \leftarrow$ extract static instruction chain from G_I
- 5: **if** $\text{closed}(C_D)$ and $\text{closed}(C_I)$ **then**
- 6: $T \leftarrow$ hybrid; $C \leftarrow C_D$
- 7: attach C_I as latent supporting evidence
- 8: $S, R \leftarrow$ score and attribute C_D with support from C_I

- 9: **else if** $\text{closed}(C_D)$ **then**
- 10: $T \leftarrow$ observed; $C \leftarrow C_D$
- 11: $S, R \leftarrow$ score and attribute C_D
- 12: **else if** $\text{closed}(C_I)$ **then**
- 13: $T \leftarrow$ instruction; $C \leftarrow C_I$
- 14: $S, R \leftarrow$ score and attribute C_I
- 15: **else**
- 16: $T \leftarrow$ no_chain; $C \leftarrow \emptyset$
- 17: $S, R \leftarrow$ weak signal severity and insufficient chain cause
- 18: **return** (S, C, R, T)

4.6. Security Considerations

ProvLoom limits dynamic execution to a sandboxed context with synthetic sensitive files, bounded resources, and controlled tool interfaces. Each sandbox is created for a review run and destroyed after the run finishes. The system does not provide real credentials or unrestricted host access during review. Static instruction analysis also stays within local artifacts, avoiding remote downloads, installer execution, or script execution when recovering latent setup and maintenance paths.

5. The Design of Benchmark

5.1. Benchmark Goals

Existing prompt injection and agent security benchmarks focus on agent tasks, injected instructions, tool misuse, and outcome labels such as attack success or data leakage [10], [22], [52], [54]. Public skill corpora and skill security studies provide useful real world samples, but their annotations usually do not expose the full evidence structure needed for chain recovery [5], [21], [47]. In particular, they often lack case level labels for source, relay, sink, trigger condition, root cause, runnability boundary, expected evidence type, and hard benign lookalikes. These gaps make it difficult to evaluate whether an analyzer recovered a reviewable security chain or only matched isolated suspicious indicators.

We design our benchmark around this missing evidence structure. Each case is a skill bundle with manually checked security semantics, fixed labels, and an expected chain when a chain exists. The benchmark includes both malicious and benign cases, covers dynamic and static instruction evidence, and includes hard benign examples that share surface patterns with risky skills. This design supports evaluation of chain recovery, root cause attribution, evidence type assignment, and resistance to over escalation.

5.2. The Construction of Benchmark

Each benchmark case contains a skill bundle and a manifest. The skill bundle is the artifact under review, and the manifest defines the evaluation oracle. The manifest records the case identifier, family, label, evaluation status, trigger type, expected behavior, primary chain, root cause, and rationale. When a case belongs to a lookalike group, the manifest also records the group identifier and pair role.

The manifest represents each expected chain with a source, zero or more relays, and a sink. Source annotations distinguish sensitive local files, generated artifacts, public files, and low sensitivity files. Relay annotations describe whether the flow is direct or passes through files, tool calls, or multiple steps. Sink annotations distinguish external endpoints, approved endpoints, local sinks, and upload destinations. Trigger annotations record whether the behavior is activated by direct tool use, command construction, LLM mediated choice, or a mixed trigger.

All labels and expected chains are fixed before baseline evaluation. The generator stores expected endpoints, edges, and root causes directly in the manifest. Manual checking verifies that each manifest matches the intended security semantics of the generated skill case. The benchmark oracle remains separate from ProvLoom outputs and baseline analyzer outputs.

The benchmark contains 147 cases: 104 malicious cases and 43 benign cases. The labels are assigned through manual expert review of the generated skill bundles and manifests. The malicious cases cover six attack families, and the benign cases cover local workflows, policy compliant outward workflows, defensive examples, and malicious lookalikes. Appendix B provides the detailed family distribution and representative abbreviated `SKILL.md` excerpts. The full skill bundle and manifest define the complete case oracle.

The benign cases are not clean controls. Benchmark constructs 43 lookalikes to create false positive pressure: 29 hard benign local workflows and 14 policy benign outward workflows. These cases reuse suspicious looking terms such as audit, report, export, mirror, note, and inventory, and some preserve staged relays or transfer structure. The benign semantics differ from malicious cases: the workflow either terminates locally, uses public or low sensitivity data, or sends data only to an approved sink. Benchmark documents 16 malicious and benign lookalike pairs that preserve similar wording, relay shape, or outward facing workflow surface while changing the security semantics through the source, relay path, or sink policy.

5.3. Evaluation Setup

The benchmark records evaluation status separately from the malicious or benign label. 24 cases fall outside automatic dynamic replay. 12 cases are design level cases for unsafe actions induced by the LLM, where the risky action depends on an interactive model decision. Another 12 cases are partially stubbed mixed flow cases whose security semantics are represented in manifests but whose dynamic replay would require unavailable external conditions or nondeterministic agent behavior. The remaining 123 cases form the dynamic runnable group and complete under bounded execution.

All benchmark runs are executed locally on a CPU workstation with a 13th Gen Intel(R) Core(TM) i7-13700H processor at 2.40 GHz. Cases that require LLM mediated action selection use the DeepSeek API. The internal evaluation uses four analysis modes. Static analyzes the skill bundle without runtime telemetry. Rule applies behavior

rules to completed executions and reports behavior labels. Rule+EPG adds execution provenance graph construction and primary chain recovery. Rule+EPG+Filtering removes irrelevant runtime nodes from the exported chain after recovery. The evaluation records the decision, recovered primary chain, root cause, evidence type, latency, and skip reason when a case does not run dynamically.

The evaluation runs static screening baselines on the skill directories. Static baselines are evaluated according to their output contracts. A scanner that reports indicators or severity levels is compared on detection and false positive behavior, not on chain metrics that require an exported source, relay, and sink chain. Chain metrics are computed only for systems that produce a chain object.

Controlled Evaluation. All benchmark experiments run in controlled local environments with synthetic files, bounded resources, and no real user credentials. The experiments do not execute against third party accounts or production systems, and they do not create observable security impact outside the evaluation environment.

5.4. Ground Truth and Metrics

The manifest defines the ground truth for each case. Each manifest declares the primary chain, behavior class, root cause, dynamic runnability status, and rationale. Labels and expected chains are fixed before evaluation and are independent of ProvLoom outputs and baseline outputs. The benchmark uses source, relay, and sink as the abstraction for comparing chains with manifest annotations. Provenance based investigation systems motivate this chain centered evaluation, where the quality of the causal path matters in addition to the alert decision [12], [16], [17].

Edge F1 compares the semantic edge set in the recovered chain with the semantic edge set in the manifest chain. We treat the recovered edges as predictions and the manifest edges as ground truth. Edge precision is the fraction of recovered edges that also appear in the manifest chain, and edge recall is the fraction of manifest edges recovered by the analyzer. Edge F1 is the harmonic mean of these two values. When no chain is recovered, Edge F1 is set to zero for malicious cases with a manifest chain.

Endpoint accuracy measures whether the recovered source and sink match the manifest. Complete chain rate measures whether the recovered source, ordered relay list, and sink all match the manifest. It is stricter than endpoint accuracy because the recovered relay structure must also match the expected chain.

Partial chain usefulness gives credit when the recovered chain contains a meaningful fragment but misses the complete manifest chain. A useful partial chain must contain at least one correct endpoint or one correct semantic edge. The score is computed as a binary case score and averaged across evaluated cases. This metric separates empty behavior labels from incomplete but reviewable chain fragments.

Root cause accuracy measures whether the analyzer assigns the root cause label. Detection rate measures the fraction of malicious cases that receive a malicious or high

risk decision within the evaluation boundary. False positive rate measures the fraction of benign cases that receive the same decision. Latency measures wall clock analysis time per case and is reported as an average.

The metrics are grouped by purpose. Detection rate and false positive rate measure behavior. Endpoint accuracy, Edge F1, complete chain rate, and partial chain usefulness measure chain recovery. Root cause accuracy measures attribution quality. Latency measures analysis cost.

5.5. Ablation Variants

We use modular ablations to evaluate which parts of ProvLoom contribute to skill review performance. Each variant runs on the same benchmark cases and is evaluated against the same manifest oracle. The oracle contains the expected label, evidence type, source, optional relays, sink, root cause, and rationale.

Rule. The Rule variant is a rule based detector over dynamic observations and local artifacts. It flags suspicious indicators such as file access, network activity, command patterns, URLs, and sensitive terms. This variant does not build an execution provenance graph and does not require a connected source to sink path.

Rule+EPG. The Rule+EPG variant adds an execution provenance graph to the rule based detector. The graph represents observed processes, files, tool calls, LLM decisions, and network actions as typed vertices and edges. This variant recovers dynamic chains by connecting observed reads, writes, tool calls, process events, and network actions into source, relay, and sink paths.

Rule+EPG+Filtering. The Rule+EPG+Filtering variant adds path filtering after dynamic chain recovery. The filtering step removes or downranks loader activity, interpreter state, temporary files, environment reads, and other nodes that do not explain the source to sink connection. This variant tests whether noise suppression improves chain precision and benign case discrimination.

All variants use the same inputs, labels, and scoring rules. The comparison measures chain recovery, root cause attribution, evidence assignment, and benign discrimination.

6. The Evaluation of ProvLoom

6.1. Benchmark Results

ProvLoom effectiveness. Table 1 reports the benchmark results. In the full ProvLoom configuration, shown as Rule+EPG+Filtering in the table, the dynamic evaluation completes 123 of the 147 benchmark cases. The completed set contains 80 malicious cases and 43 benign cases, while 24 cases fall outside the bounded dynamic runnability boundary. Within the completed set, ProvLoom detects all 80 malicious cases and does not escalate any of the 43 benign cases, giving a detection rate of 1.00 and a false positive rate of 0.00.

The chain recovery results show that the detected cases are supported by reviewable evidence paths. ProvLoom

reaches 0.9765 endpoint accuracy, 0.9338 Edge F1, 0.8676 complete chain rate, and 0.9765 partial chain usefulness. It also reaches 0.9775 root cause accuracy, showing that the recovered chains usually support the intended attribution. The average latency is 7233.04 ms per completed case, which reflects the cost of bounded execution and postprocessing.

The latency results show the cost of bounded execution. Static analysis finishes in 4.46 ms on average. Dynamic baselines take about 6.7 to 7.2 seconds per completed case. Rule+EPG has similar average latency to Rule, which indicates that execution time dominates the cost in this benchmark. Filtering adds a small amount of postprocessing time and does not change the reported decision.

Ablation study. Table 2 summarizes the coverage of the four analysis modes on the same Benchmark cases. Static removes bounded execution and analyzes only the skill bundle. Rule keeps bounded execution but removes the execution provenance graph and primary chain recovery. Rule+EPG adds graph construction and primary chain recovery. Rule+EPG+Filtering removes irrelevant runtime nodes after recovery.

Static analysis is fast and reaches zero false positives on Benchmark, but Static recovers fewer complete chains than the provenance based modes. Rule detects malicious cases but cannot explain cases as source, relay, and sink chains, so Rule receives zero scores for endpoint accuracy, Edge F1, complete chain rate, and partial chain usefulness.

Benign lookalike cases create false positive pressure in Benchmark. Rule misclassifies 4 of the 14 policy benign outward workflows, giving a false positive rate of 0.2857 on this subset, and it separates only 14 of the 16 documented lookalike pairs. These errors occur in outward looking but policy compatible cases, where the surface workflow resembles unauthorized transfer but the source or sink policy is benign. In contrast, Rule+EPG+Filtering produces no false positives on all 43 benign lookalikes and correctly separates all 16 lookalike pairs.

Rule+EPG improves endpoint accuracy from 0.8261 to 0.9765, Edge F1 from 0.8261 to 0.9338, and complete chain rate from 0.6957 to 0.8676 compared with Static. Rule+EPG also suppresses the false positives under Rule by checking whether runtime events close the expected path. Filtering preserves every benchmark decision and every chain recovery score while shortening exported chains for review.

6.2. Error and Boundary Analysis

The skipped cases fall into two groups. 12 LLM mediated action cases depend on an unsafe tool choice or action argument made through an interactive model decision. Treating them as deterministic dynamic replays would make the result depend on prompt state, model sampling, or policy mediation outside the fixed benchmark run. Another 12 external tool or platform cases require credentials, user confirmation, real service state, or later maintenance actions. The sandbox does not provide these capabilities during automatic replay, so the intended behavior cannot be observed as dynamic evidence.

TABLE 1. BENCHMARK AND ABLATION RESULTS.

Method	Detection		Chain recovery				Attribution and cost	
	Detection rate	False positive rate	Endpoint accuracy	Edge F1	Complete chain rate	Partial chain usefulness	Root cause accuracy	Average latency (ms)
Static	1.00	0.00	0.8261	0.8261	0.6957	0.8261	0.4421	4.46
Rule	1.00	0.0930	0.00	0.00	0.00	0.00	0.8000	6699.93
Rule+EPG	1.00	0.00	0.9765	0.9338	0.8676	0.9765	0.9775	6687.85
Rule+EPG+Filtering (ProvLoom)	1.00	0.00	0.9765	0.9338	0.8676	0.9765	0.9775	7233.04

Note: EPG denotes execution provenance graph. Higher is better for all metrics except false positive rate and average latency.

TABLE 2. BENCHMARK COVERAGE BY BASELINE.

Baseline	Cases	Done	Skip	Mal.	Ben.
Static	147	147	0	104	43
Rule	147	123	24	80	43
Rule+EPG	147	123	24	80	43
Rule+EPG+Filtering	147	123	24	80	43

Note: Mal. denotes malicious cases; Ben. denotes benign cases.

Errors among completed cases mainly come from run-time noise and competing chain candidates. Loader files, temporary artifacts, helper calls, generated relays, and multiple outward endpoints can obscure the expected source or sink. Filtering reduces this noise, but it cannot create evidence for behavior that the bounded run does not expose. This boundary keeps dynamic metrics limited to observed behavior and leaves latent setup or maintenance risks to static instruction evidence.

6.3. Comparison with Other Skill Analysis Tools

We compare ProvLoom with public skill security analysis tools that can scan skill artifacts. This comparison evaluates screening behavior, since these tools report findings or risk levels and do not export source, relay, and sink chains. We therefore use detection counts and accuracy for this comparison, while reserving chain recovery metrics for ProvLoom and its internal variants.

We include SkillScan, Cisco AI Defense Skill Scanner, and SkillFortify as public skill security baselines [5], [8], [21]. All three tools were evaluated from GitHub snapshots downloaded on April 25, 2026. Table 3 shows that the existing tools can identify many malicious cases, but their screening outputs do not provide the evidence structure needed for chain recovery. ProvLoom addresses a different review target: it recovers a reviewable chain, assigns an evidence type, and attributes the root cause.

TABLE 3. SKILL SECURITY SCANNER RESULTS ON BENCHMARK.

System	Version	TP	FP	FN	TN	Accuracy
SkillScan	GitHub snapshot, 2026-04-25	96	22	8	21	0.796
Cisco AI Defense	GitHub snapshot, 2026-04-25	104	43	0	0	0.707
SkillFortify	GitHub snapshot, 2026-04-25	92	22	12	21	0.769
ProvLoom	Artifact snapshot, 2026-05-17	80	0	0	43	1.000

Note: TP denotes true positives. FP denotes false positives. FN denotes false negatives. TN denotes true negatives. Public scanners are evaluated on all 147 cases. ProvLoom is evaluated on the 123 dynamic runnable cases.

7. Real-world Audit of Skill Ecosystem

7.1. Public Skill Corpus and Manual Labels

We collected 88,769 public skill samples from skill markets on March 16, 2026. We used the public Cisco AI Defense Skill Scanner to prescreen the collected samples [8]. The scanner flagged 470 samples with likely security risks. To build a manual review set with both risky and safe-looking cases, we combined these 470 flagged samples with 200 randomly sampled samples that were not flagged by the scanner.

The manual audit covers these 670 public skill samples. Each sample is labeled as malicious, ambiguous, or benign, following the detailed annotation protocol in Appendix C. A malicious label requires concrete security evidence that the sample expresses malicious or dangerous intent. An ambiguous label records incomplete evidence, such as a suspicious setup path, a risky capability request, or a partial chain that does not close. A benign label records ordinary workflow, documentation, defensive tooling, local reporting, or normal setup text. Manual annotation labels 21 samples as malicious, 79 as ambiguous, and 570 as benign. The real world audit evaluates how ProvLoom supports review and triage on selected public skills, not whether these proportions reflect the natural distribution of malicious and benign skills in the ecosystem.

7.2. Audit Findings

ProvLoom escalates samples whose highest finding severity is high or critical. This policy escalates 26 samples; manual review confirms 7 malicious, 7 ambiguous, and 12 benign samples. The result shows that the high-severity queue gives reviewers a compact set with a substantial malicious fraction.

Across the audited snapshot, manual review confirms 21 malicious samples. These samples mainly abuse setup and maintenance paths. ProvLoom links the visible skill purpose to instructions that install OpenClaw agent or CLI helpers, download protected archives, execute external scripts, replace local skill content, or transfer control to helper code outside the skill package. These evidence chains explain risk in samples whose `SKILL.md` files contain little direct executable code.

We cross-check the 21 malicious samples against public advisories, public proofs of concept, and security research

reports. At least 13 samples have related public information under this broader matching policy. The matched sources describe the same behavior classes that ProvLoom reports, including setup instructions used as malware delivery wrappers, password-protected archives, OS-specific external downloads, obfuscated or remote scripts, and hidden instruction content [6], [9], [20], [25], [43], [48]. This overlap provides external validation for the audit findings.

7.3. Closed Chains in Confirmed Malicious Samples

ProvLoom recovers closed chains for 9 of the 21 confirmed malicious samples. A closed chain connects the relevant source, intermediate steps, and security impact in one reviewable path, which lets reviewers inspect why the sample was escalated. Figure 3 shows an updater style sample in which ProvLoom recovers an instruction derived chain from local setup and maintenance artifacts. In this case, ProvLoom extracts the external `openclawcli` installation instructions as the trust boundary crossing, links them to the daily `clawdbot` cron registration as the control transfer step, and follows the scheduled update commands to the impact sink. The recovered chain starts with external `openclawcli` acquisition, continues through daily `clawdbot` cron registration and scheduled execution, and reaches update authority over Clawdbot and installed skills. The example shows that ProvLoom treats updater style text as strong instruction evidence only when local artifacts support the full path from trust boundary crossing to control transfer and security impact.




Figure 3. Updater style sample and recovered instruction chain.

8. Related Work

8.1. Skill Security

Agent skills package reusable instructions, tool descriptions, helper code, setup text, dependencies, and other artifacts for later invocation [21], [55]. This packaging model exposes risks that overlap with LLM application risks and agentic application risks. OWASP lists prompt injection, sensitive information disclosure, supply chain risk, excessive agency, tool misuse, identity and privilege abuse, unexpected

code execution, memory and context poisoning, and agentic supply chain vulnerabilities as major risk categories [30], [31]. In the skill setting, these risks can appear as hidden instructions, unsafe helper scripts, broad capability requests, data exfiltration flows, fetch and execute setup steps, persistent configuration, update mechanisms, or sensitive account connection [5], [21], [47]. SkillJect studies skill based prompt injection, where poisoned skills use modified `SKILL.md` instructions and auxiliary helper scripts to steer tool execution after the skill is loaded [18]. Skill bundles also inherit software supply chain risks from malicious packages, install scripts, dependency channels, maintainer controlled distribution paths, and update mechanisms [11], [29], [34], [35], [41], [56].

Prior work proposes skill analysis and detection frameworks from several directions. SkillScan analyzes skill artifacts and reports recurring vulnerability patterns across prompt injection, data exfiltration, privilege escalation, and supply chain risks [21]. SkillFortify studies formal analysis for agentic skill supply chains, including trust modeling, capability verification, and dependency resolution [5]. Semia audits agent skills by translating interfaces and prose instructions into a structured representation and checking security properties over that representation [47]. Public scanners such as Cisco AI Defense Skill Scanner, ClawVet, and other skill scanners provide static review and ecosystem vetting for skill packages [1], [8], [37]. These tools are useful for candidate discovery, marketplace screening, and preliminary triage. Their outputs usually take the form of indicators, severity levels, static findings, or formal verdicts.

Ecosystem studies show that skill security risks are not isolated cases. SkillScan, Semia, SkillFortify, SkillJect, and Under the Hood of `SKILL.md` examine public skill corpora, skill enabled agents, or skill registry settings and report recurring risks in real skill artifacts [5], [18], [21], [34], [47]. These measurements motivate review methods that go beyond detecting suspicious text or isolated indicators. ProvLoom complements prior skill security tools by checking whether dynamic evidence or static instruction evidence closes a source to sink chain and by reporting the evidence type behind the decision.

8.2. Provenance

System provenance research uses process, file, socket, and audit events to reconstruct attack causality and explain how an observed outcome occurred. Early intrusion backtracking and provenance systems connect runtime events into causal histories for investigation [16], [19], [33]. Later systems use provenance graphs for detection, triage, query, and attack reconstruction over noisy host telemetry [7], [12], [14], [15], [24]. Other provenance based systems align attack behavior with audit records or refine traces for threat hunting and investigation [17], [26], [27], [49].

Static analysis and source to sink reasoning provide another basis for explaining security paths before execution. Static vulnerability analysis, taint analysis, value flow analysis, and code property graphs reason about how data or

control may move through code and program relations [23], [40], [42], [50]. ProvLoom uses related path reasoning, but the evidence boundary is different. Host provenance systems usually target endpoint traces. Static analysis usually targets code or program artifacts. ProvLoom targets skill review, where relevant evidence may come from bounded execution, local skill artifacts, tool calls, generated files, LLM selected actions, and setup or maintenance instructions.

9. Discussion

9.1. Lessons Learned

Skills carry a chain risk inside the agent runtime and are easy to share and reuse, but a bundle can mix prose instructions, helper scripts, setup steps, update logic, metadata, and tool descriptions. The real world audit shows that confirmed malicious samples often rely on setup and maintenance paths, such as remote script execution, archive unpacking or local skill replacement. These paths can shape future agent behavior even when a bounded run does not expose direct data exfiltration. A risky skill is not just a bad prompt or a suspicious tool call. It can become a persistent extension point that changes local state, expands tool authority, or delegates trust to external code. The community should treat skill publication and installation as a security boundary, with clearer provenance, stronger review practice, and evidence based reports that explain why a skill is risky.

9.2. Limitations

ProvLoom is designed for evidence triage and review support, not complete malicious skill detection or full host forensics. The current framework uses bounded sandbox execution and local artifact analysis, so it can miss behavior that depends on real credentials, external platform state, delayed maintenance actions, or remote content unavailable during review. The public audit is also limited by data collection: samples come from skill markets, the snapshot reflects one collection point, and prescreening may bias the reviewed set toward risks found by the prescreening scanner. Future work should expand the collection window, include more skill sources and platforms, and improve calibration with additional manually labeled cases.

9.3. Responsible Disclosure

We initiated responsible disclosure for the 21 confirmed malicious skills identified in the audit. 13 of them have been publicly confirmed and assigned vulnerability identifiers [6], [9], [20], [25], [43], [48]. The platform is currently assessing the impact of the remaining 8 samples.

10. Ethical Considerations

We strictly adhere to the ethical principles of the Menlo Report [3] and best practices for measurement research [32].

All deployments and experimental evaluations are conducted in controlled environments. The automated testing process implemented strict frequency control, and the malicious skill execution interval was set to 10 seconds to ensure that no additional payload was imposed on any infrastructure. All evaluation activities are conducted for academic research purposes only, and any discovered vulnerabilities are not exploited for malicious attacks. Our evaluation does not attack any external infrastructure, does not involve human subjects, and does not collect any user privacy information. Each skill is executed in a fresh sandbox, and the container in which the skill is run will be destroyed immediately after each execution. Agents in the sandbox are given minimal permissions and strict context isolation, only having access to the test workspace, synthetic files, and controlled tool interfaces needed for the current use case. All tool calls are audited and automated high-risk operations are prohibited, and sensitive paths involving real credential use or host modification are transferred to the manual audit process. The malicious skills built in the experiment are only used for evaluation purposes and do not contain malicious content such as harmful URLs or real vulnerability exploitation codes. The skill artifacts used in real-world auditing are collected manually from open platforms.

11. Conclusion

Skill security review needs chain explanations, not only alerts. ProvLoom separates chains observed during bounded execution from chains recovered from setup and maintenance instructions. On the 147-case Benchmark, Rule+EPG achieves 0.9765 endpoint accuracy, 0.9338 Edge F1, and a 0.8676 complete-chain rate, while behavior-only runtime rules export no primary chain. On 670 manually annotated public samples, the audit identifies 21 malicious, 79 ambiguous, and 570 benign cases. Together, the controlled benchmark and public audit show that closed-chain evidence can make agentic skill review more precise.

References

- [1] Syed Abbas. skill-scanner: Static security scanner for SKILL.md files and AI agent skill packages. <https://github.com/syedabbast/skill-scanner>, 2026. Accessed: 2026-05-13.
- [2] Hengyu An, Jinghui Zhang, Tianyu Du, Chunyi Zhou, Qingming Li, Tao Lin, and Shouling Ji. IPIGuard: A novel tool dependency graph-based defense against indirect prompt injection in LLM agents. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 1023–1039, Suzhou, China, November 2025. Association for Computational Linguistics.
- [3] Michael Bailey, David Dittrich, Erin Kenneally, and Doug Maughan. The menlo report. *IEEE Security & Privacy*, 10(2):71–75, 2012.
- [4] Adam Bates, Dave Jing Tian, Kevin R. B. Butler, and Thomas Moyer. Trustworthy whole-system provenance for the linux kernel. In *Proceedings of the 24th USENIX Security Symposium*, pages 319–334. USENIX Association, 2015.
- [5] Varun Pratap Bhardwaj. Formal analysis and supply chain security for agentic AI skills. *arXiv preprint arXiv:2603.00195*, 2026.

- [6] Geoffrey Chan. bad-humanizer: Proof-of-concept demonstrating prompt injection in AI agent skills via hidden comments and invisible unicode tag characters. <https://github.com/gchan/bad-humanizer>, 2026. GitHub repository. Accessed: 2026-05-26.
- [7] Zijun Cheng, Qiujian Lv, Jinyuan Liang, Yan Wang, Degang Sun, Thomas Pasquier, and Xueyuan Han. KAIROS: Practical intrusion detection and investigation using whole-system provenance. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2024.
- [8] Cisco AI Defense. Cisco AI Defense Skill Scanner. <https://github.com/cisco-ai-defense/skill-scanner>, 2026. Accessed: 2026-05-13.
- [9] CNVD. Security advisory on hidden malicious command execution in some agent skill packages (Skills). <https://www.cnvd.org.cn/webinfo/show/12366>, May 2026. English translation of the Chinese advisory title. Accessed: 2026-05-17.
- [10] Edoardo Debenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems* 37, 2024.
- [11] Ruian Duan, Omar Alrawi, Ranjita Pai Kasturi, Ryan Elder, Brendan Saltaformaggio, and Wenke Lee. Towards measuring supply chain attacks on package managers for interpreted languages. *arXiv preprint arXiv:2002.01139*, 2020.
- [12] Peng Gao, Xusheng Xiao, Zhichun Li, Kangkook Jee, Fengyuan Xu, Sanjeev R. Kulkarni, and Prateek Mittal. AIQL: Enabling efficient attack investigation from system monitoring data. In *Proceedings of the 2018 USENIX Annual Technical Conference*, 2018.
- [13] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. *arXiv preprint arXiv:2302.12173*, 2023.
- [14] Xueyuan Han, Thomas Pasquier, Adam Bates, James Mickens, and Margo Seltzer. UNICORN: Runtime provenance-based detector for advanced persistent threats. In *Proceedings of the Network and Distributed System Security Symposium*, 2020.
- [15] Wajih Ul Hassan, Shengjian Guo, Ding Li, Zhengzhang Chen, Kangkook Jee, Zhichun Li, and Adam Bates. NoDoze: Combatting threat alert fatigue with automated provenance triage. In *Proceedings of the Network and Distributed System Security Symposium*, 2019.
- [16] Md Nahid Hossain, Sadegh M. Milajerdi, Junao Wang, Birhanu Eshete, Rigel Gjomemo, R. Sekar, Scott Stoller, and V. N. Venkatakrishnan. SLEUTH: Real-time attack scenario reconstruction from COTS audit data. In *Proceedings of the 26th USENIX Security Symposium*, 2017.
- [17] Yang Ji, Sangho Lee, Evan Downing, Weiren Wang, Mattia Fazzini, Taesoo Kim, Alessandro Orso, and Wenke Lee. RAIN: Refinable attack investigation with on-demand inter-process information flow tracking. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017.
- [18] Xiaojun Jia, Jie Liao, Simeng Qin, Jindong Gu, Wenqi Ren, Xiaochun Cao, Yang Liu, and Philip Torr. SkillJect: Effectively automating skill-based prompt injection for skill-enabled agents. *arXiv preprint arXiv:2602.14211*, 2026.
- [19] Samuel T. King and Peter M. Chen. Backtracking intrusions. In *Proceedings of the 19th ACM Symposium on Operating Systems Principles*, 2003.
- [20] Koi Security. ClawHavoc: 341 malicious clawed skills found by the bot they were targeting. <https://www.koi.ai/blog/clawhavoc-341-malicious-clawedbot-skills-found-by-the-bot-they-were-targeting>, 2026. Accessed: 2026-05-26.
- [21] Yi Liu, Weizhe Wang, Ruitao Feng, Yao Zhang, Guangquan Xu, Gelei Deng, Yuekang Li, and Leo Zhang. Agent skills in the wild: An empirical study of security vulnerabilities at scale. *arXiv preprint arXiv:2601.10338*, 2026.
- [22] Yupei Liu, Yuqi Jia, Runpeng Geng, Jinyuan Jia, and Neil Zhenqiang Gong. Formalizing and benchmarking prompt injection attacks and defenses. In *Proceedings of the 33rd USENIX Security Symposium*, 2024.
- [23] V. Benjamin Livshits and Monica S. Lam. Finding security vulnerabilities in java applications with static analysis. In *Proceedings of the 14th USENIX Security Symposium*, 2005.
- [24] Shiqing Ma, Xiangyu Zhang, and Dongyan Xu. ProTracer: Towards practical provenance tracing by alternating between logging and tainting. In *Proceedings of the Network and Distributed System Security Symposium*, 2016.
- [25] Jason Meller. From magic to malware: How OpenClaw’s agent skills become an attack surface. <https://1password.com/blog/from-magic-to-malware-how-openclaws-agent-skills-become-an-attack-surface>, February 2026. Accessed: 2026-05-26.
- [26] Sadegh M. Milajerdi, Birhanu Eshete, Rigel Gjomemo, and V. N. Venkatakrishnan. POIROT: Aligning attack behavior with kernel audit records for cyber threat hunting. In *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, 2019.
- [27] Sadegh M. Milajerdi, Rigel Gjomemo, Birhanu Eshete, R. Sekar, and V. N. Venkatakrishnan. HOLMES: Real-time APT detection through correlation of suspicious information flows. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2019.
- [28] Luc Moreau and Paolo Missier. PROV-DM: The PROV data model. W3C Recommendation, 2013. Accessed: 2026-05-21.
- [29] Marc Ohm, Henrik Plate, Arnold Sykosch, and Michael Meier. Backstabber’s knife collection: A review of open source software supply chain attacks. In *Detection of Intrusions and Malware, and Vulnerability Assessment*, 2020.
- [30] OWASP GenAI Security Project. OWASP Top 10 for Large Language Model Applications 2025. <https://owasp.org/www-project-top-10-for-large-language-model-applications/>, 2025. Accessed: 2026-05-26.
- [31] OWASP GenAI Security Project. OWASP Top 10 for Agentic Applications 2026. <https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>, 2026. Accessed: 2026-05-26.
- [32] Craig Partridge and Mark Allman. Ethical considerations in network measurement papers. *Commun. ACM*, 59(10):58–64, September 2016.
- [33] Thomas Pasquier, Xueyuan Han, Thomas Moyer, Adam Bates, Olivier Hermant, David Eysers, Jean Bacon, and Margo Seltzer. Runtime analysis of whole-system provenance. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*, 2018.
- [34] Shoumik Saha, Kazem Faghieh, and Soheil Feizi. Under the hood of SKILL.md: Semantic supply-chain attacks on AI agent skill registry. *arXiv preprint arXiv:2605.11418*, 2026.
- [35] Justin Samuel, Nick Mathewson, Justin Cappos, and Roger Dingle-dine. Survivable key compromise in software update systems. In *Proceedings of the 17th ACM Conference on Computer and Communications Security*, 2010.
- [36] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems* 36, 2023.
- [37] Mohib Shaikh. ClawVet. <https://github.com/MohibShaikh/clawvet>, 2026. Accessed: 2026-05-13.
- [38] Yongliang Shen, Kaitao Song, Xu Tan, Dongsheng Li, Weiming Lu, and Yueting Zhuang. HuggingGPT: Solving AI tasks with ChatGPT and its friends in Hugging Face. In *Advances in Neural Information Processing Systems* 36, 2023.
- [39] Jiawen Shi, Zenghui Yuan, Guiyao Tie, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Prompt injection attack to tool selection in LLM agents. In *Proceedings of the Network and Distributed System Security Symposium*, 2026.

- [40] Yulei Sui and Jingling Xue. SVF: Interprocedural static value-flow analysis in LLVM. In *Proceedings of Compiler Construction*, 2016.
- [41] Santiago Torres-Arias, Hammad Afzali, Trishank Karthik Kuppusamy, Reza Curtmola, and Justin Capps. in-toto: Providing farm-to-table guarantees for bits and bytes. In *Proceedings of the 28th USENIX Security Symposium*, 2019.
- [42] Omer Tripp, Marco Pistoia, Stephen J. Fink, Manu Sridharan, and Omri Weisman. TAJ: Effective taint analysis of web applications. In *Proceedings of the 30th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '09*, pages 87–97. ACM, 2009.
- [43] VirusTotal. From automation to infection: How OpenClaw AI agent skills are being weaponized. <https://blog.virustotal.com/2026/02/from-automation-to-infection-how.html>, February 2026. Accessed: 2026-05-26.
- [44] Eric Wallace, Kai Xiao, Reimar Leike, Lilian Weng, Johannes Heidecke, and Alex Beutel. The instruction hierarchy: Training LLMs to prioritize privileged instructions. *arXiv preprint arXiv:2404.13208*, 2024.
- [45] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023.
- [46] Peiran Wang, Yang Liu, Yunfei Lu, Yifeng Cai, Hongbo Chen, Qingyou Yang, Jie Zhang, Jue Hong, and Ye Wu. AgentArmor: Enforcing program analysis on agent runtime trace to defend against prompt injection. *arXiv preprint arXiv:2508.01249*, 2025.
- [47] Hongbo Wen, Ying Li, Hanzhi Liu, Chaofan Shou, Yanju Chen, Yuan Tian, and Yu Feng. Semia: Auditing agent skills via constraint-guided representation synthesis. *arXiv preprint arXiv:2605.00314*, 2026.
- [48] Wunderwuzzi. Scary agent skills: Hidden unicode instructions in skills ... and how to catch them. <https://embracethered.com/blog/posts/2026/scary-agent-skills/>, February 2026. Accessed: 2026-05-26.
- [49] Yulai Xie, Dan Feng, Yuchong Hu, Yan Li, Staunton Sample, and Darrell D. E. Long. Pagoda: A hybrid approach to enable efficient real-time provenance-based intrusion detection in big data environments. *IEEE Transactions on Dependable and Secure Computing*, 17(6):1283–1296, 2020.
- [50] Fabian Yamaguchi, Nico Golde, Daniel Arp, and Konrad Rieck. Modeling and discovering vulnerabilities with code property graphs. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2014.
- [51] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.
- [52] Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
- [53] Haiyue Zhang, Yi Nian, and Yue Zhao. Agent audit: A security analysis system for LLM agent applications. *arXiv preprint arXiv:2603.22853*, 2026.
- [54] Hanrong Zhang, Jingyuan Huang, Kai Mei, Yifei Yao, Zhenting Wang, Chenlu Zhan, Hongwei Wang, and Yongfeng Zhang. Agent security bench (ASB): Formalizing and benchmarking attacks and defenses in LLM-based agents. *arXiv preprint arXiv:2410.02644*, 2024.
- [55] Yingli Zhou, Shu Wang, Yaodong Su, Wenchuan Du, Yixiang Fang, and Xuemin Lin. A comprehensive survey on agent skills: Taxonomy, techniques, and applications. *arXiv preprint arXiv:2605.07358*, 2026.
- [56] Markus Zimmermann, Cristian-Alexandru Staicu, Cam Tenny, and Michael Pradel. Small world with high risks: A study of security threats in the npm ecosystem. In *Proceedings of the 28th USENIX Security Symposium*, 2019.

Appendix A. Artifact and Reproducibility Support

The artifact package records Benchmark by individual cases rather than only through aggregate tables. The package includes full case manifests, generator templates, deterministic generation parameters or random seeds, expected source, relay, and sink annotations, root cause labels, dynamic runnability flags, skipped case rationales, baseline outputs for each case, and aggregate summary scripts. Each manifest declares the source, relay, sink, behavior classes, benign or malicious status, and root cause before any baseline is executed. Benchmark labels are not derived from ProvLoom outputs.

The individual case records support the construct validity boundary described in the main body. Reviewers can inspect the declared task, generated skill, expected chain, and recovered output as separate artifacts. The package also exposes endpoint errors, false positives, skipped cases, and chain reconstruction failures for each case. The aggregate scripts recompute the reported endpoint, edge, chain, root cause, and latency metrics from the individual case records.

Appendix B. Benchmark Family Details

Table 4 reports the benchmark family distribution and evaluation focus. Figure 4 shows representative abbreviated SKILL.md excerpts from six malicious benchmark families. Each excerpt highlights the source, relay, sink, or trigger pattern used by the benchmark oracle.

TABLE 4. BENCHMARK CASE FAMILIES.

Family	Cases	Label	Evaluation focus
Direct sensitive exfiltration	16	Mal.	Sensitive local source and external sink without an intermediate relay.
Staged or relay exfiltration	36	Mal.	Sensitive data carried through generated files or tool call relays before transfer.
Unauthorized external transfer	16	Mal.	Generated artifacts sent to upload destinations without policy approval.
Unsafe command construction	12	Mal.	Sensitive file exposure or unsafe local execution through command construction.
Unsafe action induced by the LLM	12	Mal.	Tool choice and action selection mediated by the LLM.
Mixed flow with several steps	12	Mal.	Sensitive sources, several relays, and external sinks in the same case.
Hard benign local workflows	29	Ben.	Outward looking wording with workflows that terminate locally.
Policy benign outward workflows	14	Ben.	Public data or approved uploads with surface form similar to risky transfer.

Note: Mal. denotes malicious cases; Ben. denotes benign cases.




Figure 4. Representative abbreviated SKILL.md excerpts from six malicious benchmark families.

Appendix C. Manual Audit Protocol

The goal of the audit is to assign review labels and record evidence quality for selected public skills. The audit is used for triage analysis and case validation. It is not used to estimate the natural prevalence of malicious skills in the public ecosystem.

Review Unit. The review unit is a public skill bundle. Reviewers inspect local artifacts, including SKILL.md, README files, setup documents, maintenance notes, helper scripts, metadata, configuration files, and referenced local resources when available. The review focuses on whether the bundle contains evidence of risky behavior, whether the evidence forms a closed path, and whether the path supports a malicious, ambiguous, or benign label.

Label Definitions. A sample is labeled malicious when the available evidence supports reviewer agreement that the skill expresses malicious or dangerous intent. This label requires more than a suspicious keyword, URL, or command. The evidence must show a meaningful path toward security impact, such as data exposure, unsafe execution, persistent state change, unauthorized update behavior, or sensitive account access.

A sample is labeled ambiguous when the bundle contains a risk signal but the available evidence does not support a confirmed malicious label. Ambiguous cases include partial setup paths, risky capability requests, remote resources whose contents are unavailable, instructions that require user confirmation or credentials, and cases where intent cannot be determined from local artifacts alone.

A sample is labeled benign when the evidence indicates ordinary workflow, documentation, local reporting, defensive tooling, test code, or normal setup behavior. Benign samples may still contain commands, URLs, security terms, or example attack strings. The label depends on the reviewed context and whether the evidence forms a risky closed path. **Evidence Fields.** For each reviewed sample, the annotation records the final label, reviewer notes, the primary reason for the label, and whether the evidence contains a valid chain. Chain validity is recorded separately from the final label. A valid chain can appear in a confirmed malicious sample or in an ambiguous sample. The final label also considers intent, context, and whether the evidence is complete enough to support escalation.

Boundary Cases. The audit keeps incomplete evidence below the confirmed malicious threshold. A setup instruction that points to remote content is not enough by itself when local artifacts do not show what the remote content does. A command example is not enough by itself when the surrounding text clearly places the command in documentation, testing, or defensive analysis. A sensitive capability request is not enough by itself when no path connects the capability to misuse or security impact.

Reviewer Handling. Reviewers inspect the local artifacts and record the evidence that supports the label. When a sample contains conflicting signals, reviewers prefer the weaker label unless the evidence supports a closed risky path and malicious or dangerous intent. This rule reduces over escalation of documentation, defensive examples, and ordinary setup text. Confirmed risky samples are retained for responsible disclosure handling when reliable maintainer or platform contact channels are available.

Released Records. The artifact package records the reviewed sample identifier, final human label, reviewer notes, chain validity field, and evidence category. Released records avoid credentials, live endpoints, exploit details, and other content that would make misuse easier. The released annotations are intended to support reproducibility of the audit result while preserving the ethical boundary described in Section 10.

Appendix D. Confirmed Malicious Samples

Table 5 lists the 21 samples labeled malicious in the manual audit. The sample key is the artifact identifier used in the released annotation records. The malicious category is normalized from the manual root-cause label and reviewer notes.

TABLE 5. CONFIRMED MALICIOUS SAMPLES IN THE MANUAL AUDIT.

#	Skill	Sample key	Malicious category
1	auto-updater	auto-updater-2yq87	External agent installation and persistent update control
2	updater	autoupdate	Supply-chain persistence and environment control
3	bad-humanizer	bad-humanizer	Hidden instruction and deceptive file creation
4	clawhub	c1lawhub	Supply-chain update and environment control
5	pdf	pdf-1wso5	Remote install and remote script execution
6	rlm	pi-rlm	Supply-chain installation
7	polymarket	polytrading	Supply-chain installation
8	pomodoro	pomodoro	Supply-chain installation
9	postey	posteyskills	Supply-chain installation
10	skill-security-audit	skill-security-audit	Data exfiltration
11	visualize-with-libraries	skill-visualize-r-direct	Malicious test fixture and broad permissions
12	soulsys	soulsys-skill	Supply-chain attack
13	auto-updater	update	Supply-chain attack
14	auto-updater	updater	Supply-chain attack
15	helper-tool	version-demo	Explicit compromise and backdoor instructions
16	wallet-tracker	wallet-tracker-0ghsk	Supply-chain attack
17	x-trends	x-trends-0heof	Supply-chain attack
18	yahoo-finance	yahoo-finance-1h2ji	Supply-chain attack
19	youtube-summarize	youtube-summarize-3hazy	Supply-chain attack
20	youtube-thumbnail-grabber	youtube-thumbnail-grabber-w7har	Supply-chain attack
21	youtube-video-downloader	youtube-video-downloader-5qfuw	Supply-chain attack

LLM Usage Statement

LLMs were used for editorial purposes in this manuscript, and all outputs were inspected by the authors to ensure accuracy and originality. The use was limited to wording refinement and formatting consistency. All technical claims, experimental results, tables, and figures were produced and verified by the authors.